

Forest_Fire_logistic_regression

Bowei Zhang

2023-03-27

Data Import and Preprocessing

import data

```
df_raw = read_csv('forest_fire.csv')

## Rows: 517 Columns: 14
## -- Column specification -----
## Delimiter: ","
## chr (2): month, day
## dbl (12): area, X, Y, FFMC, DMC, DC, ISI, temp, RH, wind, rain, id
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.

head(df_raw)
```

```
## # A tibble: 6 x 14
##   area      X      Y month day  FFMC  DMC   DC  ISI  temp  RH  wind  rain
##   <dbl> <dbl> <dbl> <chr> <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1  0.1      7      5 mar  fri   86.2  26.2  94.3   5.1   8.2   51   6.7   0
## 2  0.1      7      4 oct  tue   90.6  35.4  669.    6.7  18    33   0.9   0
## 3  0.1      7      4 oct  sat   90.6  43.7  687.    6.7  14.6  33   1.3   0
## 4  0.1      8      6 mar  fri   91.7  33.3  77.5    9    8.3   97    4   0.2
## 5  0.1      8      6 mar  sun   89.3  51.3  102.    9.6  11.4  99   1.8   0
## 6  0.1      8      6 aug  sun   92.3  85.3  488    14.7  22.2  29   5.4   0
## # ... with 1 more variable: id <dbl>
```

preprocess

month and day

```
months <- c("jan", "feb", "mar", "apr", "may", "jun", "jul", "aug", "sep", "oct", "nov", "dec")
days <- c("mon", "tue", "wed", "thu", "fri", "sat", "sun")

df <- df_raw %>%
```

```

mutate(month = case_when(
  month %in% months ~ match(month, months)
)) %>%
mutate(day = case_when(
  day %in% days ~ match(day, days)
))

df$month <- factor(df$month,
  levels = 1:12,
  labels = months)
df$day <- factor(df$day,
  levels = 1:7,
  labels = days)

```

```

df <- df %>%
  mutate(area_size = case_when(
    area == 0.1 ~ 0,
    TRUE ~ 1
  ))

```

model fitting

```

mod.fit.full <- logistf(area_size ~ X + Y + month + day + temp + RH + wind + rain,
  data = df,
  family = binomial)
backward(mod.fit.full, slstay = 0.20)

```

```

## Step 0 : starting model
## Step 1 : removed day (P= 0.9183893 )
## Step 2 : removed rain (P= 0.9764783 )
## Step 3 : removed RH (P= 0.7457029 )
## Step 4 : removed Y (P= 0.6493052 )
## Step 5 : removed wind (P= 0.245984 )

```

```

## logistf(formula = area_size ~ X + month + temp, data = df, family = binomial)
## Model fitted by Penalized ML
## Confidence intervals and p-values by Profile Likelihood
##
## Coefficients:
## (Intercept)          X    monthfeb    monthmar    monthapr    monthmay
## -1.98988120  0.07035215  1.32101470  0.63969417  0.99891912  1.17180159
##   monthjun    monthjul    monthaug    monthsep    monthoct    monthnov
##  0.78285793  1.16594726  1.11693453  1.31432971  0.38279110  0.09355447
##   monthdec      temp
##  4.47496944  0.03183064
##
## Likelihood ratio test=26.87089 on 13 df, p=0.01295737, n=517

```

```
mod.fit <- logistf(formula = area_size ~ X + month + temp,
                  data = df,
                  family = binomial)
summary(mod.fit)
```

```
## logistf(formula = area_size ~ X + month + temp, data = df, family = binomial)
##
## Model fitted by Penalized ML
## Coefficients:
##               coef      se(coef)   lower 0.95   upper 0.95      Chisq
## (Intercept) -1.98988120 1.55836660 -6.924144670  0.54301896 2.276357887
## X           0.07035215 0.03935827 -0.006605016  0.14806907 3.208922529
## monthfeb    1.32101470 1.61486532 -1.354502797  6.29957503 0.828619873
## monthmar    0.63969417 1.58543526 -1.962833537  5.59483003 0.183905882
## monthapr    0.99891912 1.68485842 -1.858134314  6.03417388 0.404133076
## monthmay    1.17180159 1.94398002 -2.372716038  6.44629277 0.396018798
## monthjun    0.78285793 1.65770688 -2.001064590  5.79596910 0.251709663
## monthjul    1.16594726 1.63032927 -1.547699739  6.15695175 0.615396949
## monthaug    1.11693453 1.59513918 -1.509226012  6.07979189 0.596566199
## monthsep    1.31432971 1.58684144 -1.290882559  6.27065704 0.863219278
## monthoct    0.38279110 1.66042149 -2.415817873  5.39740138 0.056476261
## monthnov    0.09355447 2.25855008 -5.387344636  5.59846562 0.001715567
## monthdec    4.47496944 2.12305277  1.051932792 10.31685482 7.068089513
## temp        0.03183064 0.02094624 -0.009071621  0.07327184 2.323734305
##
##               p method
## (Intercept) 0.131360029      2
## X           0.073237697      2
## monthfeb    0.362671911      2
## monthmar    0.668037854      2
## monthapr    0.524962447      2
## monthmay    0.529152511      2
## monthjun    0.615873815      2
## monthjul    0.432762969      2
## monthaug    0.439891184      2
## monthsep    0.352839365      2
## monthoct    0.812154626      2
## monthnov    0.966961544      2
## monthdec    0.007846888      2
## temp        0.127413820      2
##
## Method: 1-Wald, 2-Profile penalized log-likelihood, 3-None
##
## Likelihood ratio test=26.87089 on 13 df, p=0.01295737, n=517
## Wald test = 20.84932 on 13 df, p = 0.07596292
```

```
mod.fit.full2 <- logistf(formula = area_size ~ FFMC + DMC + DC + ISI,
                        data = df,
                        family = binomial)
backward(mod.fit.full2, slstay = 0.20)
```

```
## Step 0 : starting model
## Step 1 : removed ISI (P= 0.8263427 )
```

```
## Step 2 : removed DMC (P= 0.6828909 )
## Step 3 : removed FFMC (P= 0.3426695 )

## logistf(formula = area_size ~ DC, data = df, pl = FALSE, family = binomial)
## Model fitted by Penalized ML
## Confidence intervals and p-values by Wald
##
## Coefficients:
## (Intercept)          DC
## -0.3349555602  0.0007599967
##
## Likelihood ratio test=4.580487 on 1 df, p=0.03233803, n=517

mod.fit.2 <- logistf(formula = area_size ~ DC, data = df,
                     pl = FALSE,
                     family = binomial)
summary(mod.fit.2)

## logistf(formula = area_size ~ DC, data = df, pl = FALSE, family = binomial)
##
## Model fitted by Penalized ML
## Coefficients:
##               coef      se(coef)   lower 0.95 upper 0.95    Chisq
## (Intercept) -0.3349555602 0.2144776274 -7.553240e-01 0.08541286 2.438992
## DC          0.0007599967 0.0003567634 6.075318e-05 0.00145924 4.537981
##
##               p method
## (Intercept) 0.11835225      1
## DC          0.03315067      1
##
## Method: 1-Wald, 2-Profile penalized log-likelihood, 3-None
##
## Likelihood ratio test=4.580487 on 1 df, p=0.03233803, n=517
## Wald test = 5.390093 on 1 df, p = 0.02025139
```

Visualization

```
# Create fitted df
new.df <- data.frame(
  DC = with(df,
    seq(min(DC),
      max(DC),
      by = 0.1),
  )
)
new.df$fitted <- predict(mod.fit.2, newdata = new.df, type = "response")

# ggplot2
ggplot() +
  geom_line(data = new.df,
    aes(x = DC,
```

```

    y = fitted,
    group = 1)) +
  geom_jitter(data = df,
    aes(x = DC,
        y = area_size),
    height = 0.05,
    width = 0.05,
    alpha = 0.5) +
  theme_bw()

```

